## ECSE 446/546: Realistic/Advanced Image Synthesis

Assignment 2: Progressive Monte Carlo Estimation Due: Tuesday, October 24<sup>th</sup>, 2023 at 11:59pm EST on myCourses Final weight: 25%

1 Assignment Policies and Submission Process 1.1 Late Policy, Collaboration & Plagiarism, Python/Library Usage Rules 2 Pixel Anti-Aliasing and Progressive Renderer 3 Phong Normal Interpolation 4 Ambient Occlusion and Importance Sampling 4.1 Uniform Spherical Sampling 4.2 Cosine Importance Sampling 5 You're Done!

Contents

# 1 Assignment Policies and Submission Process

Download and modify the standalone Python script we provide on myCourses, renaming the file according to your student ID as YourStudentID.py

As usual, every new file you submit on myCourses will override the previous submission, and we will only grade the final submitted file.

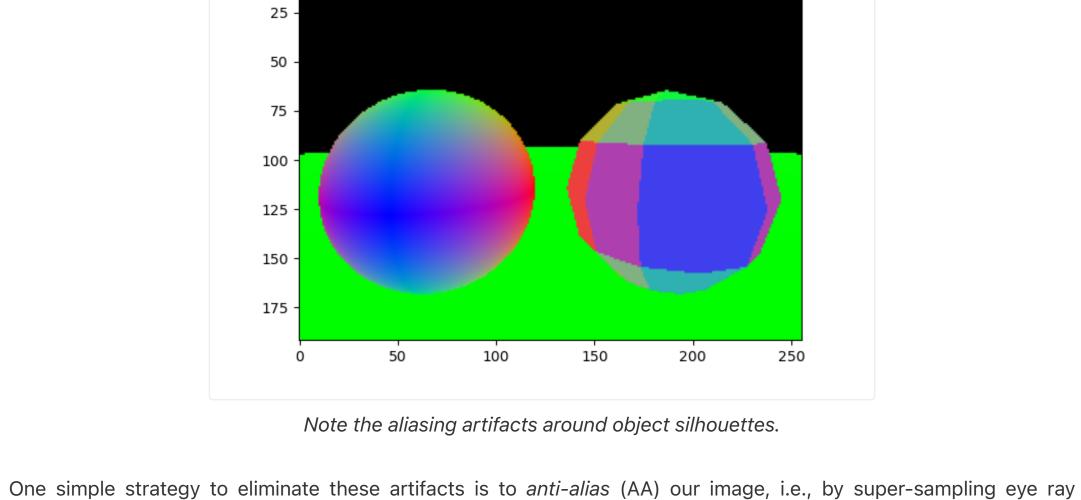
### For late policy, collaboration & plagiarism, Python language and library usage rules, please refer to the Assignment 0 handout.

1.1 Late Policy, Collaboration & Plagiarism, Python/Library Usage Rules

2 Pixel Anti-Aliasing and Progressive Renderer When generating your eye rays in Assignment 1, we exclusively sampled a viewing direction through the center of each square pixel. When the directly visible geometry varies spatially at a rate higher than our pixel grid resolution,

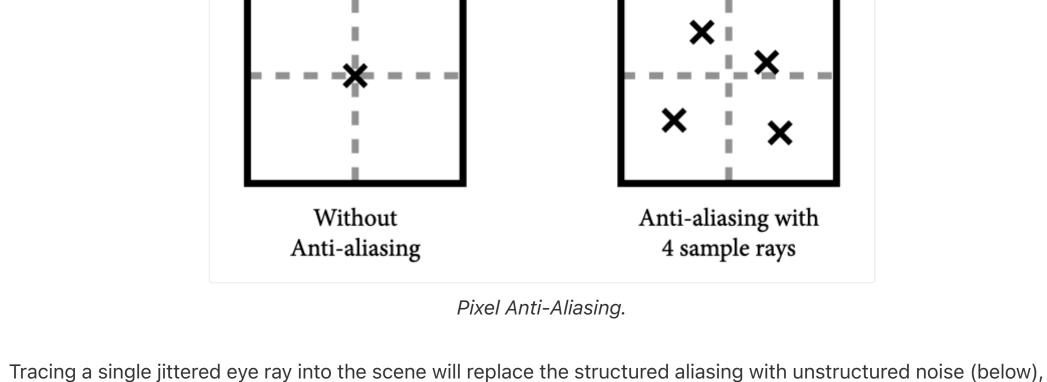
our resulting image can suffer from so-called "jaggies" — aliasing artifacts that manifest themselves primarly at the silhouettes of visible objects:

current spp: 1 of 1

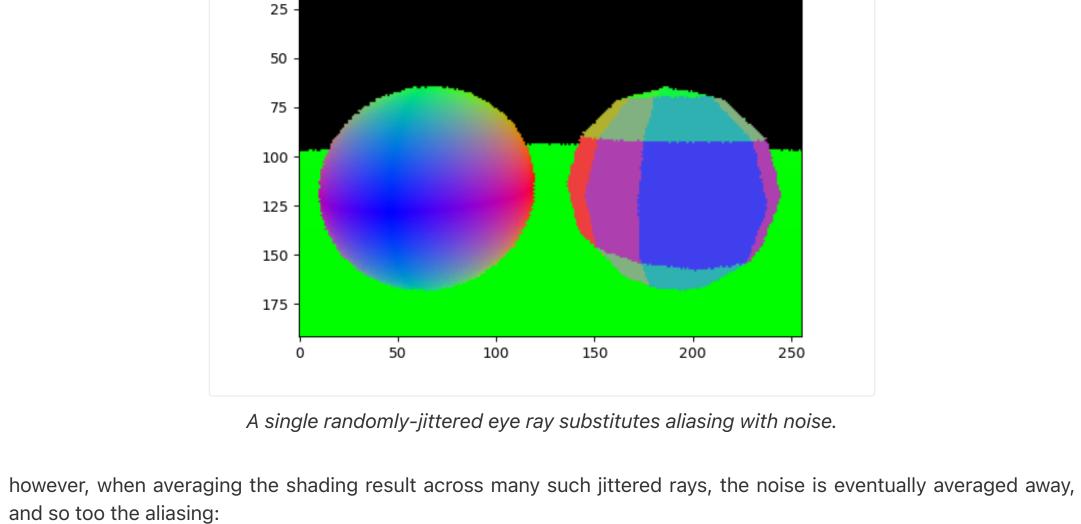


directions over each pixel's area.

Concretely, for each pixel, instead of considering a single ray through its center, we will average the contribution (e.g., the shading) across many primary rays. These jittered eye ray directions will instead be generated by picking a random location (uniformly over the area of the pixel) when generating your primary rays:

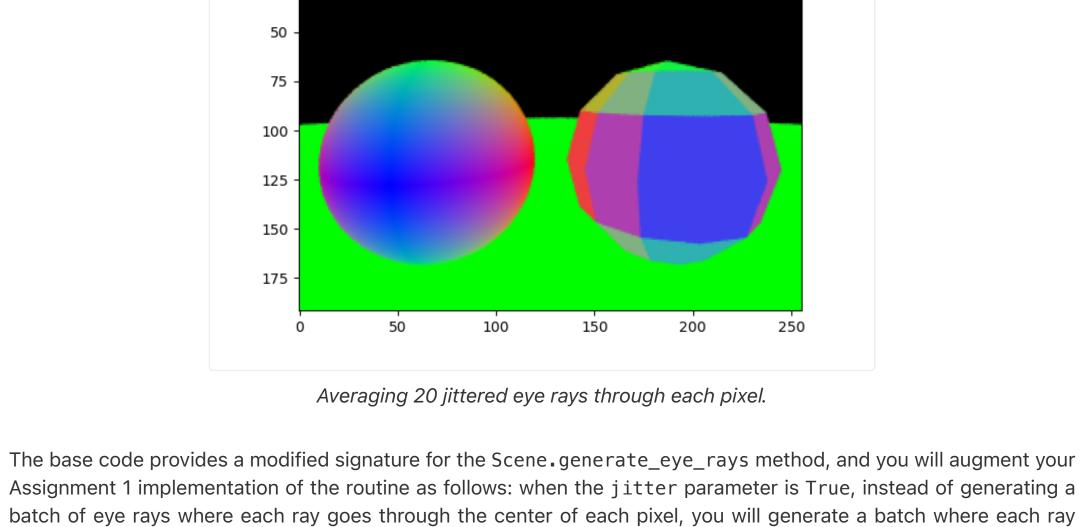


current spp: 1 of 1



current spp: 20 of 20

25 -



goes through a random location in the area of each pixel.

results of the Scene. render routine.

Scene render generates a batch of Scene width \* Scene height eye rays (with or without jittering, depending on the Scene.progressive\_render\_display's jitter parameter) and computes shading at each of these shading points. For the first two deliverables, you can use the debug implementation of Scene. render that we furnish, which simply visualizes the primary shading points' normals.

Concretely, Scene progressive\_render\_display expects the following arguments: • jitter: True if the generated eye rays should be jittered at each progressive iteration (i.e., AA enabled), False

Moreover, you will complete the implementation of a progressive accumulation renderer in

Scene.progressive\_render\_display that handles the logic of iteratively computing, averaging and displaying the

otherwise (i.e., no AA), total\_spp: the total number of samples per pixel<sup>1</sup> for the final rendered image, • spppp: the number of *Monte Carlo integration samples* to trace **per progressive rendering iteration** (for Deliverables 3 and 4),

sampling routines required in Deliverables 3 (for both ECSE 446 and 546) and 4 (only for ECSE 546). <sup>1</sup> More precisely, the final image should have np.ceil(total\_spp / sppp) samples, as in the base code. You cannot vectorize your Scene progressive render\_display routine over rendering passes (since

• sampling\_type: a parameter that will also be passed to Scene. render to select between the importance

**Additional Notable Changes in the Assignment 2 Base Code** Unlike Assignment 1, the scene can now contain both analytic sphere and discretized mesh objects; as such, we have abstracted different geometry types using the Geometry class, with Sphere and Mesh subclasses.

you need to display the result after each iteration); simply use a for loop over the individual passes.

One important distinction between how Scene intersect and Geometry intersect interact — compared to Assignment 1 — is that Geometry.intersect must now return **both** the intersection distance **and** hit point normal, for valid ray-geometry intersections. **Deliverable 1** [10 points] • augment your implementation of the Sphere.intersect from Assignment 1 to now return the hit

• augment your implementation of Scene.intersect from Assignment 1 to treat arbitrary geometry

• augment your implementation of the Scene.generate\_eye\_rays from Assignment 1 to generate a

jittered bundle of eye rays, when jitter is True (and the originally unjittered eye rays if False), • complete the implementation of Scene.progressive\_render\_display to iteratively generate eye rays, passing them to Scene. render and displaying a running average of the Scene. render output; each call of Scene. render will (eventually, in Deliverables 3 and 4) rely on sppp when computing Monte Carlo estimates of the ambient occlusion. Scene.progressive\_render\_display does not return any values. 3 Phong Normal Interpolation When shading triangle meshes, we reviewed three spatial shading strategies in class: flat, Gouraud, and Phong shading. The base code's implementation of Mesh.intersect relies on vectorized numpy to compute an ensemble of ray-

## 1. hit\_distances: an array with shape (num\_rays,) of floating point ray intersection distance parameters (or np.inf for those rays that did not intersect the mesh), 2. triangle\_hit\_ids: an integer numpy array of size (num\_rays,) with indices for the triangle in the mesh that

geometry and the Barycentric coordinates.

**Per-face Normals** 

current spp: 1 of 1

distances and hit point normals,

types, by relying on Geometry.intersect,

was intersected (and -1 for those rays that did not intersect the mesh), and 3. barys: a numpy array of size (num\_rays, 3) of the Barycentric coordinates of the ray-triangle intersection (or [0,0,0] for those rays that did not intersect the mesh).

mesh intersections. A ray bundle is tested against all triangles in the mesh and records the closest hit point. For rays

with valid intersections, Mesh.mesh\_intersect\_batch returns three numpy arrays of outputs:

The implementation that we furnish to you simply returns the triangle face normals (i.e., flat shading) for valid raymesh intersections.

You will be responsible for implementing Phong normal interpolation, relying on the intersected triangle face

More precisely, given a ray that intersects a mesh, Barycentric coordinates can be used to linearly interpolate any scalar- or vector-valued signal — defined at the vertices of the triangle — over points on the face of the triangle. For example, given a single Barycentric coordinate  $[\alpha, \beta, \gamma]$  (recalling, of course, that the vectorized output of

Mesh.mesh\_intersect\_batch provides an array of such coordinates), we can, e.g., obtain the intersection point **p** 

Of more immediate use to you, we can similarly obtain the Phong-interpolated normals for a point on the face of the

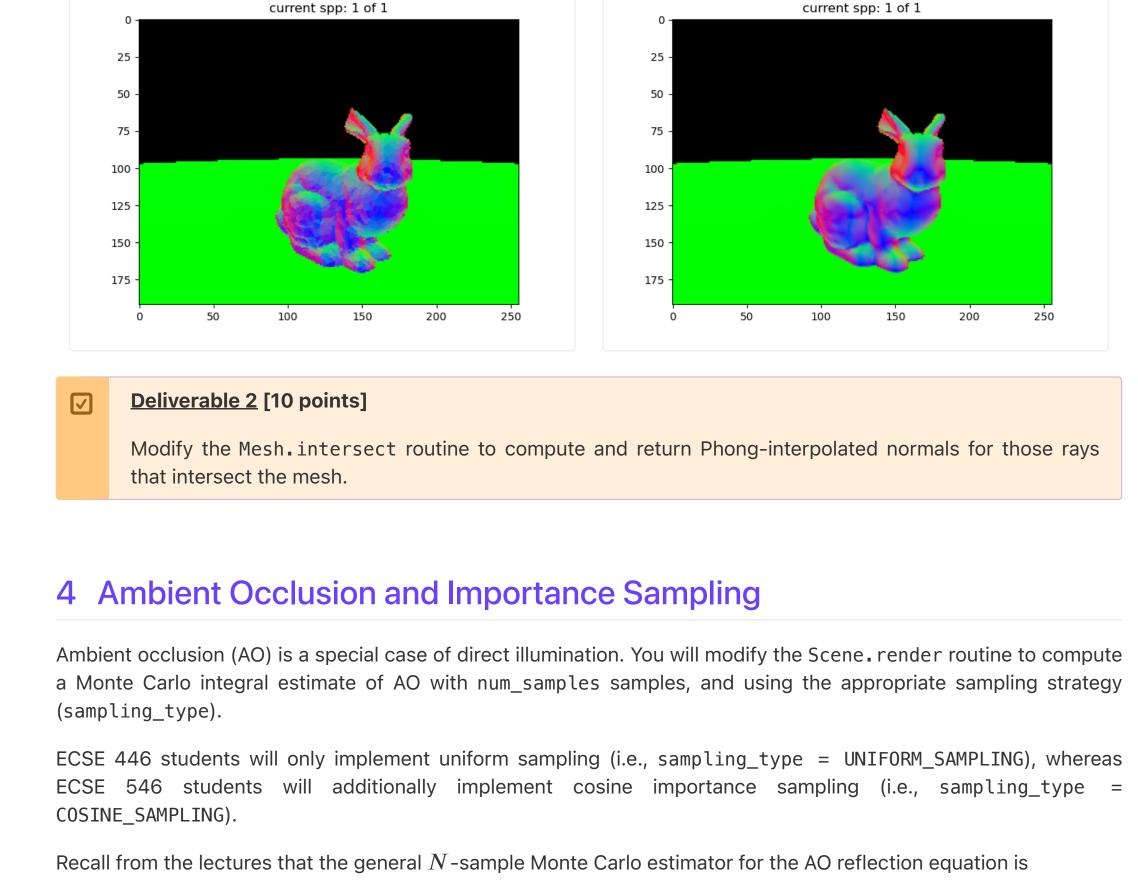
on the face of the triangle as a linear combination of the triangle's vertices, as  $\mathbf{p} = \alpha \mathbf{v}_0 + \beta \mathbf{v}_1 + \gamma \mathbf{v}_2$ .

triangle (i.e., the ray-mesh intersection point, with associated Barycentric coordinates  $[\alpha, \beta, \gamma]$ ), as  $\mathbf{n} = \alpha \, \mathbf{n}_0 + \beta \, \mathbf{n}_1 + \gamma \, \mathbf{n}_2 \,,$ where  $\mathbf{n}_i$  is the vertex normal at  $\mathbf{v}_i$ . The vertex normals are loaded at mesh initialization in the numpy array Mesh.vn.

**Phong-interpolated Per-vertex Normals** 

current spp: 1 of 1

100 -125 125 -150 150 -175 175 -200



# Geometry.brdf\_params property. 4.1 Uniform Spherical Sampling

Implement a uniform MC sampler, with your choice of hemispherical or spherical sampling (and their associated PDF

Recall from the lecture that you can generate a uniformly-sample ray direction  $\omega = (\omega_x, \omega_y, \omega_z)$  on the sphere

using two canonical random variables  $\xi_1,\xi_2$  — computed using np.random.rand — and the following

 $\omega_z = 2\xi_1 - 1 \qquad r = \sqrt{1 - \omega_z^2} \qquad \phi = 2\pi\xi_2$ 

and sampling routines), to estimate the AO integral.

each Scene. render call in the progressive rendering loop.

**Uniform AO Estimator (sphere scene)** 

current spp: 1 of 1

**Deliverable 3** [10 points]

transformation:

 $L_r(\mathbf{x}) \approx \frac{\rho}{\pi N} \sum_{j=1}^N \frac{V(\mathbf{x}, \omega_j) \max(0, \mathbf{n} \cdot \omega_j)}{p(\omega_j)} \text{ with } \omega_j \sim p(\omega),$ 

where visibility V can be evaluated by tracing a shadow ray, and the diffuse albedo ho is encoded in the

 $\omega_x = r \cos \phi$   $\omega_y = r \sin \phi$ Scene.render's sampling\_strategy parameter denotes the sampling strategy, and the number of samples N as num\_samples. In Scene progressive\_render\_display, you will pass spppp (samples per pixel per pass) as the num\_samples for

UNIFORM\_SAMPLING scenario: compute a num\_samples-sample Monte Carlo estimator of the AO reflection equation using uniform sampling. • Ensure that your Scene.progressive\_render\_display routine properly updates the render view, displaying progressively improving integral estimates (i.e., with progressively-increasing total sample counts.)

**Uniform AO Estimator (bunny scene)** 

current spp: 1 of 1

• Complete the implementation of the Scene.render function to support the sampling\_type ==

current spp: 100 of 100 current spp: 100 of 100

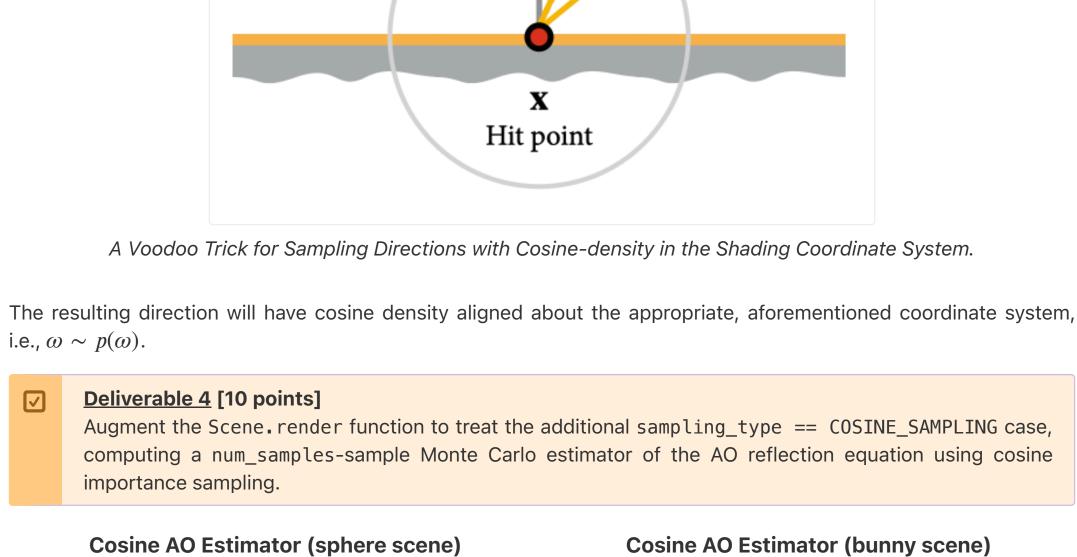
# 4.2 Cosine Importance Sampling In addition to the uniform Monte Carlo estimator, ECSE 546 students will implement cosine importance sampling.

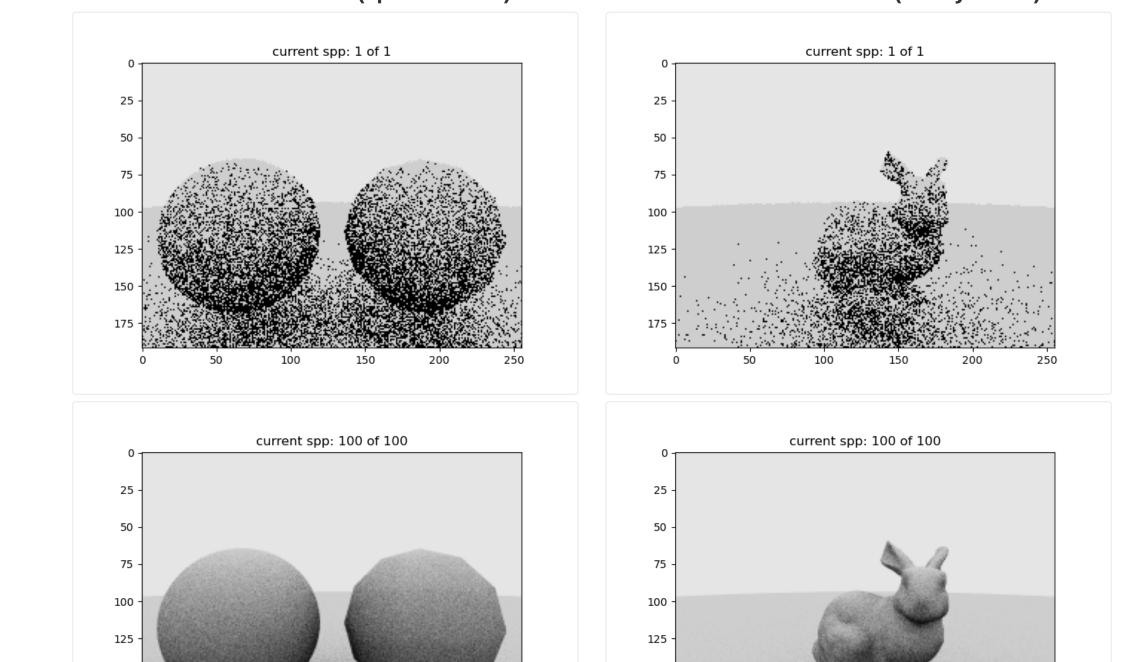
**ECSE 546 Students Only** 

# Recall that cosine importance sampling can be more effective than uniform sampling, in the AO setting. Here, the

sampling PDF is  $p(\omega) = (\cos \theta) / \pi$ , and the specialized Monte Carlo estimator simplifies expression simplifies to

 $L_r(\mathbf{x}) \approx \frac{\rho}{N} \sum_{j=1}^N V(\mathbf{x}, \omega_j) \text{ with } \omega_j \sim p(\omega).$ There are many strategies for sampling cosine-distributed points (about the coordinate system defined by the surface normal  $\bf n$  at  $\bf x$ ). Perhaps the simplest such approach is to first generate a uniform spherical sample ( $\omega_{temp}$ ), then add it to the surface normal, before finally normalizing the result as  $\omega$ :





175

formatted by <u>Markdeep 1.16</u>